

The Weight of the Adjoint

Which Symplectic Form a Neural Operator Preserves, and Whether It Survives the Lift

Sourav Sahu
Microsoft
sahusourav@microsoft.com

Gareth O’Brien
Microsoft
O.Gareth@microsoft.com

Abstract

Structure-preserving neural operators are built on a question that is usually left unasked: symplectic *with respect to which inner product*? On a non-uniform grid the discrete L^2 adjoint is $K^* = W^{-1}K^\top W$, where W collects the quadrature weights, and the choice of W silently selects which symplectic form a gradient shear preserves. We make that choice explicit, build both alternatives, and measure them. Using the symplectic defect $\|WA - A^\top W\|_F / \|WA\|_F$ computed from the autograd Jacobian, we find a clean diagonal across 1-D, 2-D and 3-D: each construction is exact (2×10^{-16}) in the form its adjoint was built from and $O(1)$ away from the other. Preserving *a* symplectic form is therefore cheap; the content lies entirely in preserving the *right* one. Because our benchmarks are discretised on uniform grids, this yields a controlled test, and the matched-form operator wins 16 of 18 paired comparisons. When we remove the lift and projection layers so that the *deployed map*—not merely its interior blocks—is symplectic, the margin widens to 14–22 $\times$. The same construction then produces the study’s strongest result: over a 2×10^4 step rollout every operator we tested diverges to relative errors of 10^5 – 10^{12} , *except* the lift-free symplectic ones, which stay bounded near 1 with energy drift of order unity. Structure preservation is not a cost, as a prior version of this study concluded; it is a benefit that is forfeited by the conventional lift–process–project wrapper and invisible unless the inner product is named. We also report a 783-run, nine-equation, three-seed benchmark, and retract an earlier symplecticity theorem of our own that conflated unit Jacobian determinant with preservation of ω .

1 Introduction

Neural operators learn maps between function spaces and have become standard surrogates for parametric PDEs (Kovachki et al., 2023; Li et al., 2021; Lu et al., 2021). A fast-growing branch argues that such operators should respect the *geometric structure* of the dynamics they emulate—symplecticity for Hamiltonian systems, or the metriplectic split when dissipation is present. The argument has a strong pedigree in finite dimensions (Greydanus et al., 2019; Jin et al., 2020; Cranmer et al., 2020; Zhong et al., 2020), and in 2026 several groups carried it into function space (Makara et al., 2026; Sulskis and Ravi, 2026; Zhang et al., 2026; Varghese et al., 2026).

Two things are almost never done. First, whether a trained network is actually structure-preserving is argued on paper rather than measured. Second—and this is our starting point—the claim is stated without naming the inner product it is relative to. On a uniform grid that omission is harmless, because the quadrature weight is a single scalar. On a non-uniform grid it is not: the weights vary by a factor of order N (Figure 1), the adjoint is $K^* = W^{-1}K^\top W$ rather than $K^\top$, and two implementations that differ only in which W they assume preserve two *different* symplectic forms.

This gives an unusually clean experiment. Build both, verify that each is exact in its own form, then ask which one helps. Our contributions:

1. **An instrument** (§5). We measure the symplectic defect directly from the autograd Jacobian, against *both* candidate inner products, scaling to 3-D with Hutchinson probes. It first falsified a symplecticity claim in our own earlier work (Appendix A).
2. **A construction** (§4). SA-CHEB: exact symplecticity on a non-periodic Chebyshev grid in 1-D, 2-D and 3-D, for any kernel and any pointwise nonlinearity (Theorem 5). We describe it and the CKINO operator it grew out of component by component.
3. **The form result** (§7.1). Each variant is exact in its own form and $O(1)$ in the other. The variant whose form matches the grid wins 16/18 paired comparisons; with the lift removed, 4/4 by 14–22×.
4. **The lift result** (§7.3). Only operators that are symplectic *end to end* survive a 2×10^4 -step rollout. A symplectic core wrapped in pointwise lift, projection and a residual update does not.
5. **A benchmark** (§6): 783 runs over nine equations in 1-D/2-D/3-D at matched capacity, three seeds, plus an unconstrained-capacity sweep showing the ranking is budget-dependent.

A correction. An earlier version of this study reported that exact symplecticity was a measurable *cost*. That conclusion was an artefact of scoring both variants against a single inner product. We state the error plainly here because it is the kind the instrument exists to catch.

2 Related work

Neural operators. The Fourier Neural Operator parameterises the integral kernel by truncated Fourier multipliers (Li et al., 2021), inheriting discretisation invariance and $O(N \log N)$ cost; DeepONet factorises the operator into branch and trunk networks (Lu et al., 2021); the general framework and universality results are due to Kovachki et al. (2023). Tensorised (Kossaifi et al., 2023) and U-Net-augmented (Wen et al., 2022) variants improve parameter efficiency. A practitioner-oriented account of FNO’s assumptions is given by Duruisseaux et al. (2026), with reference implementations in Kossaifi et al. (2024). Physics-informed variants add PDE residuals to the loss (Raissi et al., 2019; Li et al., 2023); these are *soft* constraints with no algebraic guarantee.

Structure preservation in finite dimensions. Hamiltonian (Greydanus et al., 2019) and Lagrangian (Cranmer et al., 2020) neural networks learn the generating function; SympNets (Jin et al., 2020) and Symplectic ODE-Net (Zhong et al., 2020) parameterise symplectic maps directly. These act on $\mathbb{R}^{2d}$ at a fixed discretisation and are not operators in the function-space sense.

Structure-preserving neural operators. Makara et al. (2026) introduce the Symplectic Neural Operator (SNO), composing symplectic shear blocks whose update fields are gradients of learned energy functionals, realised with self-adjoint Fourier layers on the torus. Sulskis and Ravi (2026) embed the GENERIC degeneracy structure (Grmela and Öttinger, 1997; Öttinger, 2005) into an FNO, extending finite-dimensional bracket learning (Zhang et al., 2022). Zhang et al. (2026) compose structure-preserving blocks on boundary-adapted spectral representations, and Varghese et al. (2026) apply a conditional symplectic operator to optimal control.

Where we differ. All of these are built on the torus or on a uniform grid, where $W \propto I$ and the adjoint question does not arise—the correct adjoint is obtained *by accident*. [Makara et al. \(2026\)](#) state explicitly that their formulation covers strong canonical symplectic Hilbert spaces and that extensions require further machinery. We work in the setting where the choice is forced, show that it is consequential, and measure rather than assert the outcome. Auxiliary ingredients we use—Lie-equivariant lifting ([Finzi et al., 2020](#)) and FiLM conditioning ([Perez et al., 2018](#))—are standard and not claimed as contributions.

3 Background: symplectic structure on a non-uniform grid

3.1 Notation

Table 1 collects every symbol used in the paper. Vectors are columns, $^\top$ is the plain matrix transpose, and $*$ is reserved for the adjoint with respect to a weighted inner product—the distinction between the two is the subject of the paper.

Table 1: Symbols and notation.

symbol	reads as	meaning
Ω_d	—	spatial domain, taken as $[-1, 1]$ in 1-D
d	—	number of spatial dimensions, $d \in \{1, 2, 3\}$
N	—	grid index; a 1-D grid carries $N+1$ nodes
$x_j = \cos(j\pi/N)$	—	Chebyshev–Gauss–Lobatto (CGL) nodes
$w_j > 0$	—	quadrature weight at node j
$W = \text{diag}(w_j)$	—	weight matrix; <i>diagonal and positive</i>
$W_{\text{cheb}}, W_{\text{unif}}$	—	the two candidate weight matrices: Clenshaw–Curtis weights on CGL nodes, and constant weights (a uniform grid)
$\langle f, g \rangle_W = f^\top W g$	“ f dot g in W ”	discrete L^2 inner product
K	—	a linear (kernel-integral) operator
$K^\top$	“ K transpose”	matrix transpose
$K^* = W^{-1} K^\top W$	“ K adjoint”	adjoint <i>with respect to</i> $\langle \cdot, \cdot \rangle_W$
q, p	—	the two halves of a canonical phase-space state
$z = (q, p)$	—	a point of phase space $\mathbb{P}$
ω_W	“omega”	symplectic form induced by W , eq. (2)
Ω	—	matrix of ω_W , eq. (2)
Φ	“Phi”	one shear map of phase space
$D\Phi$	—	Jacobian of Φ
F, G	—	the update (shear) fields
$A = DF$	—	Jacobian of the shear field
$R, \rho = R'$	“rho”	scalar potential and its derivative, applied pointwise
$D = \text{diag}(\rho'(\cdot))$	—	diagonal matrix of pointwise slopes
$\mathcal{E}$	“script E”	learned energy functional
∇_W	“grad in W ”	gradient with respect to $\langle \cdot, \cdot \rangle_W$
φ_r, ψ_r	“phi, psi”	the r -th left/right basis functions of the kernel
M_r	—	channel-mixing matrix of mode r
R (rank)	—	Mercer rank; context distinguishes it from the potential R
σ_r	“sigma”	learned positive Mercer scale
C, H	—	physical and hidden channel counts
$\mathcal{D}_W$	“defect”	relative symplectic defect, eq. (10)
$\ \cdot\ _F$	“Frobenius norm”	$\ X\ _F^2 = \sum_{ij} X_{ij}^2$
$[S, M] = SM - MS$	“commutator”	used in Proposition 6

3.2 The discrete inner product

Let $\Omega_d = [-1, 1]$ be discretised on the $N+1$ CGL nodes $x_j = \cos(j\pi/N)$. Functions are represented by nodal values, and the physical L^2 product is approximated by Clenshaw–Curtis quadrature (Trefethen, 2000),

$$\langle f, g \rangle_W = \int_{\Omega_d} f g dx \approx f^\top W g, \quad W = \text{diag}(w_0, \dots, w_N), \quad w_j > 0. \quad (1)$$

Phase space is $\mathbb{P} = \mathcal{H}_q \oplus \mathcal{H}_p$ with $\mathcal{H} = (\mathbb{R}^{N+1}, \langle \cdot, \cdot \rangle_W)$, carrying

$$\omega_W(z, z') = \langle q, p' \rangle_W - \langle p, q' \rangle_W = z^\top \Omega z', \quad \Omega = \begin{pmatrix} 0 & W \\ -W & 0 \end{pmatrix}. \quad (2)$$

This is the finite-dimensional shadow of the weak symplectic form on a Hilbert phase space (Weinstein, 1969; Marsden and Ratiu, 1999); the point of what follows is that the discretisation chooses W , and that choice is not innocuous.




Figure 1: The geometric fact the paper rests on. **(a)** CGL nodes cluster quadratically at the boundary; a uniform grid does not. **(b)** The Clenshaw–Curtis weights therefore vary by $78\times$ across the domain at $N = 24$, so $W \neq cI$ and consequently $K^* \neq K^\top$. On a uniform grid the two coincide, which is why a Fourier-parameterised operator never encounters the distinction. **(c)** The Chebyshev basis assumes no periodicity at the domain ends.

Definition 1 (Weighted adjoint). The adjoint of K with respect to (1) is the unique K^* with $\langle Kf, g \rangle_W = \langle f, K^*g \rangle_W$ for all f, g . Expanding, $f^\top K^\top Wg = f^\top WK^*g$, hence

$$K^* = W^{-1}K^\top W \quad (3)$$

Proposition 2 (Shear criterion). Let $\Phi(q, p) = (q, p + F(q))$ with $A = DF(q)$. Then

$$(D\Phi)^\top \Omega (D\Phi) - \Omega = \begin{pmatrix} WA - A^\top W & 0 \\ 0 & 0 \end{pmatrix}, \quad (4)$$

so Φ preserves ω_W iff A is self-adjoint in W , i.e. $WA = A^\top W$. The transposed shear behaves analogously.

Proof. $D\Phi = \begin{pmatrix} I & 0 \\ A & I \end{pmatrix}$, so $\Omega D\Phi = \begin{pmatrix} WA & W \\ -W & 0 \end{pmatrix}$ and $(D\Phi)^\top \Omega D\Phi = \begin{pmatrix} I & A^\top \\ 0 & I \end{pmatrix} \begin{pmatrix} WA & W \\ -W & 0 \end{pmatrix} = \begin{pmatrix} WA - A^\top W & W \\ -W & 0 \end{pmatrix}$. Subtract Ω . $\square$

Remark 3 (Volume preservation is not symplecticity). $\det D\Phi = 1$ for *every* shear, whatever A . Unit Jacobian determinant is Liouville volume preservation, strictly weaker than preservation of ω_W , which by Proposition 2 additionally requires $WA = A^\top W$. Conflating the two is, in our experience, the most common way an architecture described as symplectic turns out not to be; we exhibit an instance of our own in §7.1 and Appendix A.

Remark 4 (The question is under-specified). Because (4) depends on W , “is this block symplectic?” has no answer until W is named. Two implementations differing only in the W used inside the adjoint are *both* exactly symplectic—in different forms. This is the observation the rest of the paper develops.

4 Architectures

We describe two operators. CKINO is a Chebyshev kernel-integral operator that motivated this study and serves as our non-periodic baseline; SA-CHEB is the self-adjoint construction that makes the symplectic form explicit and controllable.

4.1 CKINO: a Chebyshev kernel-integral operator

CKINO replaces FNO’s Fourier multipliers with a low-rank kernel expanded in *Chebyshev coefficients*, so that no periodicity is assumed and the operator can be evaluated on any CGL grid. Figure 2 shows the pipeline; the components are as follows.




Figure 2: CKINO. The kernel is stored as Chebyshev *coefficients* and evaluated on whatever CGL grid is requested, so the parameters are independent of resolution. The (q, p) block mimics a Störmer–Verlet half-step. Its Jacobian is triangular with unit diagonal, hence volume-preserving—but, as §7.1 shows by measurement, *not* symplectic in either candidate form.

(i) **Lie-equivariant lifting.** A pointwise or short-stencil map raising the C physical fields to H hidden channels, built from generators of a small Lie group so that the lift commutes with translations of the input (Finzi et al., 2020). Two variants matter later: a *convolutional* lift (resolution-dependent stencil) and a *spectral* lift (FFT-derivative, exactly resolution-invariant). §7.8 shows these trade invariance against accuracy.

(ii) **Separable Chebyshev kernel integral.** The operator core,

$$(Kv)(x) = \sum_{r=1}^R \left(\prod_{a=1}^d \varphi_r^a(x_a) \right) \sigma_r M_r \langle \prod_a \psi_r^a, v \rangle_W, \quad (5)$$

with φ_r^a, ψ_r^a stored as $n_{\text{train}}+1$ Chebyshev coefficients per axis, $M_r \in \mathbb{R}^{H \times H}$ mixing channels and $\sigma_r > 0$ a learned Mercer scale. Cost is $O(RN^d)$ rather than $O(N^{2d})$, and because the basis functions are polynomials the same weights evaluate on any grid.

(iii) **Störmer–Verlet-style block.** Hidden channels are split into (q, p) halves and updated by alternating half-steps with a learned step size. This is where the original symplecticity claim was made, and where it fails: (5) uses *independent* φ_r and ψ_r and an unconstrained M_r , so nothing forces $k(x, y) = k(y, x)^\top$ and the Jacobian is not self-adjoint in any of the forms we consider.

(iv) **Projection.** A pointwise map back to C output channels.

4.2 SA-Cheb: making the form explicit

SA-CHEB repairs (iii) by construction. Let K be any bounded operator, $R \in C^1(\mathbb{R})$ with $\rho = R'$, and define the energy functional and its W -gradient

$$\mathcal{E}(q) = \int_{\Omega} R((Kq)(x)) dx, \quad F(q) = \nabla_W \mathcal{E}(q) = K^* \rho(Kq), \quad (6)$$

with K^* the weighted adjoint (3). The block alternates

$$\Phi^\uparrow(q, p) = (q, p + F_\theta(q)), \quad \Phi^\downarrow(q, p) = (q + G_\theta(p), p). \quad (7)$$




Figure 3: The SA-CHEB gradient shear. The only design choice is which W enters K^* : the Clenshaw–Curtis weights (W_{cheb}) or constant weights (W_{unif}). Both give an *exactly* symplectic block—in different forms. That choice, not the presence of symplecticity, is what the experiments vary.

Theorem 5 (Exact symplecticity, any dimension). *For any parameters of K , any $R \in C^1$, any resolution N and any spatial dimension $d \in \{1, 2, 3\}$, the map $\Phi^\uparrow$ of (7) satisfies $(\Phi^\uparrow)^* \omega_W = \omega_W$ exactly, and likewise $\Phi^\downarrow$; hence so does any finite alternating composition.*

Proof. Write $D = \text{diag}(\rho'((Kq)_j))$, so $A = DF = K^* DK$. Using (3),

$$WA = W(W^{-1} K^T W) DK = K^T W DK, \quad A^T W = K^T D (K^*)^T W = K^T D W K W^{-1} W = K^T D W K.$$

W and D are diagonal and therefore commute, so $WA = A^T W$; apply Proposition 2. In d dimensions the weight is the outer product $W = \bigotimes_a W_a$, still diagonal, so the argument is unchanged. $\square$

Theorem 5 says the block is symplectic in *the form its adjoint was built from*. It says nothing about any other form, and the next result shows that the omission is not a technicality: it is the whole content of the paper, and it predicts the pattern measured in Figure 5.

Proposition 6 (Which form, and only which form). *Fix a positive diagonal W_1 and build the shear field with the adjoint taken in W_1 , i.e. $F(q) = K_{W_1}^* \rho(Kq)$ with $K_{W_1}^* = W_1^{-1} K^T W_1$, so that $A = K_{W_1}^* DK$. Let W_2 be any other positive diagonal weight. Then*

$$W_2 A - A^T W_2 = [S, M] = SM - MS, \quad S := W_2 W_1^{-1}, \quad M := K^T W_1 DK, \quad (8)$$

where S is diagonal and M is symmetric. Consequently

$$(W_2 A - A^T W_2)_{ij} = (s_i - s_j) M_{ij}, \quad (9)$$

and the shear is symplectic in ω_{W_2} for every K and ρ if and only if $W_2 \propto W_1$.

Proof. Since W_1 is symmetric, $(K_{W_1}^*)^\top = (W_1^{-1}K^\top W_1)^\top = W_1 K W_1^{-1}$. Diagonal matrices commute, so $W_1 D = D W_1$ and $W_1^{-1} W_2 = W_2 W_1^{-1} = S$. Then

$$W_2 A = W_2 W_1^{-1} K^\top W_1 D K = S (K^\top W_1 D K) = S M,$$

$$A^\top W_2 = K^\top D (K_{W_1}^*)^\top W_2 = K^\top D W_1 K W_1^{-1} W_2 = (K^\top W_1 D K) S = M S,$$

which gives (8). M is symmetric because $M^\top = K^\top (W_1 D)^\top K = K^\top D W_1 K = K^\top W_1 D K = M$. Writing $S = \text{diag}(s_i)$ gives $(SM - MS)_{ij} = s_i M_{ij} - M_{ij} s_j$, which is (9).

For the equivalence: if $W_2 = c W_1$ then $S = c I$, every $s_i - s_j$ vanishes and so does the defect—recovering Theorem 5 at $c = 1$. Conversely, suppose $W_2 \not\propto W_1$, so there are indices $i \neq j$ with $s_i \neq s_j$. Choose $D = I$ (take $\rho' \equiv 1$) and $K = e_i e_i^\top + e_j e_j^\top$; then $M = K^\top W_1 K$ has $M_{ij} = 0$, so pick instead $K = (e_i + e_j) e_i^\top + (e_i + e_j) e_j^\top$, for which $M_{ij} = (w_{1,i} + w_{1,j}) \neq 0$. By (9) the (i, j) entry of the defect is $(s_i - s_j) M_{ij} \neq 0$, so the shear is not symplectic in ω_{W_2} for that parameter choice. $\square$

Remark 7 (The defect scales with the spread of W_2/W_1). Equation (9) says the mismatch is governed entirely by the *spread* of the ratio $s_i = w_{2,i}/w_{1,i}$: a larger dynamic range in $W_2 W_1^{-1}$ produces a larger defect. This is a falsifiable prediction, and Figure 5 bears it out—the off-diagonal entries grow from ≈ 0.7 in 1-D to ≈ 0.95 in 2-D and 1.0–1.3 in 3-D, precisely because the tensor-product weight $\otimes_a W_a$ has a wider dynamic range as d grows.

Low-rank realisation. With $Kv = \sum_r \varphi_r M_r \langle \psi_r, v \rangle_W$ the weighted adjoint is $K^* g = \sum_r \psi_r M_r^\top \langle \varphi_r, g \rangle_W$ (Appendix B): *swap* $\varphi \leftrightarrow \psi$ and *transpose* M . No inverse of W is ever formed and the cost stays $O(RN^d)$.

4.3 Lifted versus lift-free, and why it matters

Theorem 5 is a statement about the *block*. A deployed operator usually wraps blocks in a pointwise lift and projection and is driven residually, $x_{t+1} = x_t + \mathcal{N}(x_t)$. None of those three maps is symplectic, so the composition is not either (Figure 4). We therefore build both:

- **Lifted** (`sacheb`, `sacheb_naive`): $H \gg C$ hidden channels, maximal capacity, symplectic *core* only.
- **Lift-free** (`sacheb_pure`, `sacheb_pure_naive`): the shears act directly on the PDE’s own (q, p) fields, $H = C$, no lift, no projection, non-residual update. The *deployed map* is exactly symplectic; capacity must come from rank and depth instead. This requires a canonical (q, p) pair and so applies only to two-field problems.




Figure 4: Where the guarantee is lost. The conventional wrapper surrounds a symplectic core with three maps that are not symplectic; the composition therefore carries no guarantee. §7.3 shows this is not a technicality: at 2×10^4 steps the lifted variants diverge while the lift-free ones stay bounded.

Table 2: The operator variants compared, and the question each isolates.

variant	symplectic in	role
sacheb	W_{cheb} (core only)	Chebyshev-form gradient shear, lifted
sacheb_naive	W_{unif} (core only)	identical but uniform-form: isolates <i>which</i> form
sacheb_pure	W_{cheb} (end to end)	removes the lift: isolates <i>where</i> the guarantee holds
sacheb_pure_naive	W_{unif} (end to end)	both controls at once; the matched-form, fully-guaranteed operator
CKINO	neither	the non-gradient kernel; volume-preserving only
CKINO-strict	neither	exactly-constrained variant of CKINO
CKINO-nosymp	n/a	structure removed: isolates the block’s contribution
SNO (Makara et al., 2026)	W_{unif} (core only)	Fourier symplectic reference
FNO, T-FNO, U-FNO, U-Net, DeepONet, Transformer	n/a	unstructured references
GENERIC-FNO (Sulskis and Ravi, 2026)	metriplectic	dissipative reference (see §9)

4.4 The variants, and why each exists

Table 2 lists the comparison. The design is a 2×2 factorial—{Chebyshev, Fourier} basis $\times$ {symplectic, not}, realised as SA-CHEB/CKINO/SNO/FNO—crossed with two further controls that vary *which* form is preserved and *whether the deployed map* preserves it.

5 Measuring structure: the symplectic defect

Proposition 2 turns the question into a number. For a shear field F we form $A = DF$ by reverse-mode differentiation and report

$$\mathcal{D}_W(F) = \frac{\|WA - A^\top W\|_F}{\|WA\|_F}, \quad (10)$$

zero iff the block preserves ω_W . Crucially we evaluate (10) for *both* $W \in \{W_{\text{cheb}}, W_{\text{unif}}\}$; reporting one column is what produced the erroneous conclusion of our earlier draft.

Scaling to 3-D. A dense Jacobian costs one backward pass per degree of freedom—fine at $N = 32$ in 1-D, hopeless at 16^3 . Proposition 8 replaces it with a probe estimator whose cost is independent of the grid.

Proposition 8 (Unbiased probe estimator for the defect). *Let $M_* := WA - A^\top W$ and let $v \sim \mathcal{N}(0, I)$. Then*

$$M_* v = W(Av) - A^\top(Wv), \quad \mathbb{E}_v[\|M_* v\|^2] = \|M_*\|_F^2, \quad (11)$$

so averaging $\|M_ v\|^2$ over m independent probes is an unbiased estimator of $\|M_*\|_F^2$. Each probe costs one Jacobian-vector product (Av) and one vector-Jacobian product ($A^\top u$), both $O(1)$ in the number of degrees of freedom.*

Proof. The first identity is linearity: $(WA - A^\top W)v = W(Av) - A^\top(Wv)$, and the two terms are exactly a JVP and a VJP of the shear field, so neither requires A to be formed. For the second, $\mathbb{E}[vv^\top] = I$ gives

$$\mathbb{E}_v[\|M_*v\|^2] = \mathbb{E}_v[v^\top M_*^\top M_* v] = \text{tr}(M_*^\top M_* \mathbb{E}[vv^\top]) = \text{tr}(M_*^\top M_*) = \|M_*\|_F^2.$$

The same argument applied to WA estimates the denominator of (10), so the ratio is formed from two estimates sharing the JVP. $\square$

We use $m = 96$ probes. Against the dense path the estimator agrees to $\leq 7\%$ on $O(1)$ values and correctly returns machine zero; the dense path is retained wherever the state has at most 1200 degrees of freedom, and is what the estimator is validated against.

6 Experimental protocol

Problems. Nine PDEs: advection, heat, `wave1d`, `wave1d_dir`, Burgers, KdV, `wave2d` (32^2), `wave3d` (16^3) and `ns2d` (64^2 vorticity Navier–Stokes). All but `wave1d_dir` are periodic.

`wave1d_dir` is the discriminating rung: the wave equation on $[0, L]$ with homogeneous *Dirichlet* boundaries, uniform grid, centred differences, symplectic leap-frog, and initial conditions formed from a truncated Chebyshev series times the envelope $x(L - x)$. This mirrors the data-generation protocol of Makara et al. (2026), so the comparison is made on the concurrent method’s own terms. It is simultaneously Hamiltonian (symplecticity is meaningful) and non-periodic (a Fourier parameterisation is mismatched). The reference solver conserves energy to 2.3×10^{-4} over 500 steps and holds the boundary at exactly zero.

Design. Families are crossed with prediction mode {recursive, non-recursive `seq2seq`}, training noise {none, 2%} and an optional PDE-residual loss, giving 33 configurations in 1-D and 21 in 2-D/3-D, all matched to $\approx 25\text{k}$ parameters. Three seeds on a single hardware (NVIDIA Tesla T4): **783 training runs.**

Metrics. Relative RMS at rollout checkpoints, plus amplitude ratio, pattern correlation, spectral error and invariant drift, collapsed into a *verdict* and a *usable horizon* (last step with correlation ≥ 0.9 and amplitude ratio in $[0.7, 1.4]$). A model can post a middling RMS while being physically dead, and only the decomposition separates the two.

7 Results

7.1 Which form does each construction preserve?

Figure 5 is the central measurement. Three conclusions follow.

(i) Theorem 5 holds numerically: $2\text{--}4 \times 10^{-16}$, independent of resolution and dimension. (ii) The two variants are *both* exactly symplectic, in different forms; neither is “broken”. The off-diagonal grows with dimension (≈ 0.7 in 1-D, ≈ 0.95 in 2-D, 1.0–1.3 in 3-D) because the tensor-product weight varies more. (iii) The CKINO kernel is ≈ 1.4 in both forms, exactly the conflation flagged in Remark 3.




Figure 5: Symplectic defect of every construction against *both* inner products, over 14 grid configurations in 1-D, 2-D and 3-D, mean of three seeds. The green diagonal is the result: each gradient shear is exact in the form its adjoint was built from and $O(1)$ in the other. The CKINO kernel is ≈ 1.4 in both—volume-preserving, not symplectic.

7.2 Does the choice of form matter?

Every benchmark here is discretised on a *uniform* grid, so W_{unif} is the physically relevant form and W_{cheb} is not. Figure 6 compares models that are identical in architecture, parameter count and training recipe and differ only in the adjoint.




Figure 6: Matched-form versus mismatched-form at identical capacity (3 seeds). (a) Lifted pair: the matched-form model wins every problem, by 1.1–8.0 $\times$. (b) Lift-free pair, where the deployed map really is symplectic: the margin widens to 14–22 $\times$. Removing the non-symplectic wrapper does not just help—it amplifies the effect of getting the form right by an order of magnitude.

Over the full matrix the matched-form model wins **16 of 18** paired comparisons (the two exceptions are a tie at the divergence clip and a 1.04 $\times$ gap inside one standard deviation). For the lift-free pair it is 4/4.

Table 3: Lift-free pair: both models exactly symplectic *end to end*, differing only in which form. Relative RMS, 3 seeds.

problem	purecheb (W_{cheb})	pureunif (W_{unif})	ratio
wave1d	1.5182	0.08809	17.2×
wave1d_dir	0.1020	0.08242	1.24×
wave2d	0.9332	0.04316	21.6×
wave3d	1.0273	0.07238	14.2×

7.3 Long-horizon behaviour: the lift is decisive




Figure 7: Rollout to 2×10^4 steps, mean of 3 seeds. **Top:** relative error; the thin line marks error = signal. **Bottom:** energy drift on the model’s own trajectory. Every lifted or unstructured operator—FNO, SNO, CKINO, and the lifted SA-CHEB variants—leaves the solution manifold and reaches 10^5 – 10^{12} . The lift-free operators stay bounded. The matched-form lift-free model (green) is bounded on all four problems; the mismatched-form one (red) survives three of four.

This is the study’s strongest result and it was not visible at the few-hundred-step horizon of the main matrix. Figure 7 and Table 4 show that exact *end-to-end* symplecticity delivers precisely what the classical theory promises—bounded energy error over long times—and that nothing else in the comparison does.

Remark 9 (Read the slope with the error). A drift exponent near zero is *not* by itself evidence of health. FNO scores -0.02 on **wave1d** while sitting at relative RMS 7×10^5 : once the state has blown up, the drift ratio saturates and its slope becomes meaningless.

Independent corroboration. Sulskis and Ravi (2026) report the same qualitative effect for the metriplectic case: over 200-step rollouts, every unconstrained model they test diverges or collapses while their structure-preserving operator stays bounded. That two constructions enforcing different

Table 4: State at step 2×10^4 (mean of 3 seeds). Bold rows are bounded; the rest have left the solution manifold entirely.

problem	family	relative RMS	energy drift $ \Delta E/E $
wave1d	pureunif (lift-free, W_{unif})	0.868	4.18
	purecheb (lift-free, W_{cheb})	1.902	37.9
	SNO / FNO / lifted SA-CHEB	$2.6\text{--}7.2 \times 10^5$	$10^{11}\text{--}10^{12}$
wave1d_dir	pureunif	0.908	5.24
	purecheb	1.659	36.2
	SNO / FNO	$3.4\text{--}9.4 \times 10^5$	$10^{11}\text{--}10^{13}$
wave2d	pureunif	1.008	2.52
	purecheb	1.739	3.94
	FNO	7.3×10^5	1.7×10^{12}
wave3d	pureunif	1.047	1.19
	SNO / FNO / lifted	$2.8\text{--}8.3 \times 10^5$	$10^{11}\text{--}10^{12}$

structures (symplectic here, GENERIC there) both buy boundedness and little else suggests the effect is a property of exact structure preservation rather than of either particular bracket.

The same result seen as fields. Figures 8 and 9 show the rollout as the state itself rather than as summary curves. They make two things visible that Figure 7 cannot. First, the failure modes are *qualitatively different*: blow-up (the lifted variants and FNO), amplitude collapse (SNO), and graceful roughening at preserved amplitude (lift-free). Second, amplitude collapse produces a field that is numerically well behaved—no overflow, bounded norm—while carrying none of the physics, which is precisely the case an RMS-only evaluation at short horizon would pass.

7.4 Basis versus boundary

The converse holds too, which is what makes this a rule rather than a cherry-pick: in one dimension the winning basis follows the boundary conditions *exactly*—Fourier operators take all four periodic problems and the Chebyshev family takes the one non-periodic problem—while the Chebyshev operator additionally takes all three multi-dimensional problems. No structural prior is needed to explain the split.

7.5 Accuracy across the full matrix

Against the parameter-matched references the margins in higher dimensions are $19.8\times$ over FNO on **wave2d**, $8.8\times$ over FNO and $4.8\times$ over SNO on **wave3d**, and $1.5\times/1.4\times$ on **ns2d**. We stress that SNO is now competitive in 2-D ($1.0\times$), which the earlier version of this study could not have detected because SNO had no N -dimensional implementation there.

Table 5 lists a single best model per problem because that is the conventional presentation, but it overstates how decisive most of those rows are. Applying the tie rule of §7.7—a gap smaller than the sum of the two seed standard deviations is not a result—only five of the nine problems have a clear winner (**wave1d_dir**, **wave3d**, **ns2d** to our construction; KdV to SNO; Burgers to T-FNO). Advection, heat, **wave1d** and **wave2d** are ties. We give the resulting tally in §8.

wave1d-Dir, autoregressive rollout to 2×10^4 steps: reference solver (grey) vs operator prediction




Figure 8: wave1d_dir time lapse over the 2×10^4 -step rollout: reference solver (grey) against operator prediction. Row labels carry each family’s population error over all test trajectories, so a single plotted trajectory cannot be over-read. The *lifted* SA-Cheb and FNO blow up; SNO collapses to a low-amplitude oscillation—bounded, but carrying none of the physics; only the *lift-free* operators track the wave to 2×10^4 steps.

wave-2D displacement field: reference solver against operators over a 2×10^4 -step rollout




Figure 9: wave2d displacement field over the same rollout. The reference solver sustains a coherent interference pattern throughout. The lift-free matched-form operator degrades gracefully into noise while remaining bounded; the *lifted* variant—same shear algebra, same parameters, only the wrapper differs—has diverged by 2×10^4 , and FNO by 10^3 .




Figure 10: `wave1d_dir`, the Hamiltonian non-periodic rung. Every Chebyshev operator beats every Fourier operator: 3.2 $\times$ over the best FNO and 6.5 $\times$ over SNO, on a benchmark built to match SNO’s own data protocol. The mechanism is the Gibbs penalty the Dirichlet envelope imposes on a periodic basis.

Table 5: Best configuration per problem (relative RMS, mean $\pm$ std, 3 seeds).

problem	best model	relative RMS	basis
advection [‡]	SNO (seq2seq)	0.001196 $\pm$ 0.00016	Fourier
heat	T-FNO (recursive)	0.001090 $\pm$ 0.00031	Fourier
wave1d	SNO (seq2seq)	0.001103 $\pm$ 0.00039	Fourier
wave1d_dir	sacheb_naive (seq2seq)	0.025939 $\pm$ 0.0021	Chebyshev
Burgers	T-FNO (recursive)	0.003963 $\pm$ 0.00098	Fourier
KdV	SNO (seq2seq)	0.000511 $\pm$ 0.000064	Fourier
wave2d	sacheb_naive (seq2seq)	0.007771 $\pm$ 0.0015	Chebyshev
wave3d	pureunif (recursive)	0.072375 $\pm$ 0.0089	Chebyshev
ns2d	sacheb_naive (seq2seq)	0.174500 $\pm$ 0.0014	Chebyshev

[‡] advection is a tie within seed noise: the contending Chebyshev model varies by 15.7 $\times$ across seeds there, and the winner changed between two executions of the study (§7.7). We do not claim it.




Figure 11: Relative RMS (log scale, mean $\pm$ std over 3 seeds) for every configuration, one panel per equation, all matched to $\approx 25k$ parameters.

Accuracy vs capacity once the matched-parameter constraint is removed. Faint x marks settings that diverged.




Figure 12: Accuracy against parameter count with the matched budget removed; the dashed line marks the 25k budget used in Table 5. The ranking is budget-dependent, so the matched table should be read as “best at 25k” rather than “best architecture”.

7.6 Unconstrained capacity

Two observations. On `wave1d_dir` the two end-to-end symplectic models take the top two slots outright (0.0742 and 0.0755)—the only problem where they win the unconstrained comparison, and the only Hamiltonian non-periodic one. And cost-efficiency diverges sharply from peak accuracy: on `wave3d` SNO’s 0.00271 costs 1.36M parameters while `sacheb_naive` reaches 0.00345 with 49,750—**27× cheaper** for **1.3× worse**.

7.7 Replication: a standard we think this literature needs

Operator-learning papers are almost never replicated, including by their own authors. Seed counts of three are common, margins of a few percent are reported as wins, and there is usually no way for a reader to tell which of those wins would survive re-running the study. We are in a position to check, because the entire matrix—all 783 runs—was executed twice on the cluster. The duplication was not planned, but treating it as a control is the most useful thing we can do with it, and we argue the practice should be routine rather than accidental.

Table 6: Every headline claim under a full re-execution of the study.

claim	execution A	execution B	verdict
form ablation (lifted)	$W_{\text{unif}} 16 / W_{\text{cheb}} 1$	$W_{\text{unif}} 16 / W_{\text{cheb}} 2$	robust
lift-free ratios	17.2/1.24/21.6/14.2	identical	robust
long-horizon boundedness	lift-free only	lift-free only	robust
winner, 8 of 9 problems	—	—	stable
winner, advection	CKINO 0.001092	SNO 0.001196	<i>flipped</i>

The one claim that did not survive is instructive. On advection the contending Chebyshev configuration varies by $15.7\times$ across seeds (0.00062 to 0.00968) against SNO’s $1.4\times$. That is not a close race between two good models; it is one unstable configuration occasionally landing well. Under a single execution it would have been reported as a win for our own operator. We report advection as a tie.

Two things follow. First, as a methodological rule: **with three seeds, a difference smaller than the seed spread is not a result**. Every claim we advance in §1 survived re-execution with margins of $1.4\text{--}22\times$; the single claim that did not had a margin of $1.1\times$. Second, as a reporting standard: the seed spread should be published alongside the mean, because it is what tells a reader whether a margin is load-bearing. We report both throughout.

7.8 Supporting axes

Darcy (elliptic, Dirichlet). CKINO attains $8.7\times$ lower relative L^2 error than FNO overall and $18.3\times$ lower *at the boundary*, transferring zero-shot from 64^2 to 128^2 . This is the clearest non-periodic win for the Chebyshev basis.

Capacity ceilings. A 6k–400k sweep shows CKINO’s Burgers failure is *representational*, flat at ≈ 0.25 across a $60\times$ range, and its recursive KdV instability persists at every budget. Neither is fixed by scale.

Discretisation invariance. Training at $N = 64$ and evaluating at $N = 128$, **SNO and T-FNO win this axis outright**: they are essentially exactly invariant ($0.99\text{--}1.00$) *and* accurate, which is the combination that matters. CKINO with a convolutional lift degrades by $57\times$ on advection; the spectral lift restores invariance (0.998) but costs two to three orders of magnitude of accuracy at the training grid ($8.2 \times 10^{-4} \rightarrow 0.475$). There is no setting of our construction that is both invariant and accurate, and we do not have a fix. If a workflow requires evaluating at a resolution it did not train on, that is a decisive argument for the Fourier family and against ours.

Inference speed. Operators beat the reference solver only where the solver is expensive: FNO is $3.2\times$ on Burgers, $7.7\times$ on KdV and $5.8\times$ on **ns2d**; CKINO is $2.3\times$ on KdV but $0.07\times$ on advection.

8 Where each construction wins

We want to be explicit about the shape of this table rather than let it be read selectively. Counting a problem as a tie whenever the gap between the two best families is smaller than the sum of their seed standard deviations—the rule §7.7 argues for—only *five* of the nine problems have a clear winner: **wave1d_dir**, **wave3d** and **ns2d** to our construction, KdV to SNO, Burgers to T-FNO. The remaining four (advection, heat, **wave1d**, **wave2d**) are ties. Splitting ties evenly the tally is ours 5.0, SNO 2.5, T-FNO 1.5.

Our construction wins the two axes it was designed for—behaviour under long rollout, and accuracy on non-periodic and multi-dimensional domains. It loses the two it was not. **SNO wins discretisation invariance outright**, being simultaneously invariant and accurate where our operator can be only one or the other, and the Fourier family wins or ties every periodic 1-D problem. A reader whose workload is periodic, or who must evaluate at a resolution the model was not trained on, should use a Fourier operator.

Table 7: Honest summary. No single operator dominates; the axes separate, and two of the four axes go to the Fourier family.

construction	wins	loses
pureunif (lift-free, W_{unif})	long horizon (only bounded operator at 2×10^4 steps); wave3d at 25k; top of the unconstrained wave1d_dir sweep	needs a canonical (q, p) pair, so undefined on scalar fields; limited capacity per parameter
sacheb_naive (lifted, W_{unif})	wave1d_dir , wave2d , ns2d at 25k; best accuracy-per-parameter in 3-D	no long-horizon guarantee; diverges like the rest at 10^4 steps
CKINO	advection/heat unconstrained by 10^2 – $10^3 \times$; Darcy by $8.7 \times$ ($18.3 \times$ at the boundary)	not symplectic in any form; Burgers ceiling; KdV instability; worst discretisation invariance in the study
SNO (Makara et al., 2026)	KdV outright, and ties advection, wave1d and wave2d ; discretisation invariance outright ; competitive in 2-D	Fourier basis mismatched on non-periodic domains ($6.5 \times$ on wave1d_dir); expensive in 3-D (1.36M parameters for its best)
T-FNO (Kossaifi et al., 2023)	heat, Burgers; invariance	no structure; long-horizon divergence

9 Limitations

Re-implementations. SNO and GENERIC-FNO are our implementations from the published descriptions and were not tuned to the authors’ protocols.

Our GENERIC-FNO bug. It diverged in all 12 entries. Parameterising the PSD multiplier as $M = b^2$ with b initialised at zero places it at a saddle ($\partial M / \partial b = 2b = 0$), so the dissipative channel never activates. This is a defect of our implementation, *not* evidence about the published method, and no conclusion should be drawn from those rows.

Canonical structure only. The lift-free construction needs a canonical (q, p) pair and is therefore undefined on the scalar-field problems (advection, heat, Burgers, KdV, **ns2d**). KdV is Hamiltonian only under the non-canonical Gardner bracket, which we do not cover.

Scale and horizon. Domains are small (32 – 64 points in 1-D, 32^2 , 16^3 , 64^2), adequate to resolve the reported orderings but not a deployment benchmark. 2×10^4 steps separates bounded from unbounded but is still far short of the regimes where backward-error analysis is usually invoked.

Single hardware. All 783 runs used one accelerator type (Tesla T4) with TF32 disabled, so arithmetic is consistent across the matrix; results may differ on accelerators with different defaults.

10 Conclusion

We set out to test whether geometric structure helps a learned operator and found that the question, as usually posed, is incomplete. Symplecticity is defined relative to an inner product; on a non-uniform grid two reasonable implementations preserve two different forms, both exactly, and only measurement distinguishes them. Once the form is named, the answer is clean: matching it to the grid wins 16 of 18 comparisons, and 4 of 4 by an order of magnitude more when the deployed map—not merely its core—carries the guarantee. That same construction is the only one in our study to survive a 2×10^4 -step rollout. Structure preservation pays; it is simply forfeited by the

wrapper the field conventionally puts around it, and invisible unless you say with respect to what. We also retract a symplecticity theorem of our own that mistook unit Jacobian determinant for preservation of ω , and we think the general lesson is that such claims should be measured rather than argued—the measurement costs one line of autograd.

References

- M. Cranmer, S. Greydanus, S. Hoyer, P. Battaglia, D. Spergel, and S. Ho. Lagrangian Neural Networks. *ICLR 2020 Workshop on Integration of Deep Neural Models and Differential Equations*, 2020. arXiv:2003.04630.
- V. Duruisseaux, J. Kossaifi, and A. Anandkumar. Fourier Neural Operators Explained: A Practical Perspective. arXiv:2512.01421, 2026.
- M. Finzi, S. Stanton, P. Izmailov, and A. G. Wilson. Generalizing Convolutional Neural Networks for Equivariance to Lie Groups on Arbitrary Continuous Data. *ICML*, 2020. arXiv:2002.12880.
- S. Greydanus, M. Dzamba, and J. Yosinski. Hamiltonian Neural Networks. *NeurIPS*, 2019. arXiv:1906.01563.
- M. Grmela and H. C. Öttinger. Dynamics and thermodynamics of complex fluids. I. Development of a general formalism. *Physical Review E*, 56(6):6620–6632, 1997.
- E. Hairer, C. Lubich, and G. Wanner. *Geometric Numerical Integration: Structure-Preserving Algorithms for Ordinary Differential Equations*. Springer Series in Computational Mathematics, Vol. 31, 2nd edition, Springer, 2006.
- P. Jin, Z. Zhang, A. Zhu, Y. Tang, and G. E. Karniadakis. SympNets: Intrinsic structure-preserving symplectic networks for identifying Hamiltonian systems. *Neural Networks*, 132:166–179, 2020. arXiv:2001.03750.
- J. Kossaifi, N. Kovachki, K. Azizzadenesheli, and A. Anandkumar. Multi-Grid Tensorized Fourier Neural Operator for High-Resolution PDEs. arXiv:2310.00120, 2023.
- J. Kossaifi, N. Kovachki, Z. Li, D. Pitt, M. Liu-Schiaffini, R. J. George, B. Bonev, K. Azizzadenesheli, J. Berner, V. Duruisseaux, and A. Anandkumar. A Library for Learning Neural Operators. arXiv:2412.10354, 2024.
- N. Kovachki, Z. Li, B. Liu, K. Azizzadenesheli, K. Bhattacharya, A. Stuart, and A. Anandkumar. Neural Operator: Learning Maps Between Function Spaces with Applications to PDEs. *Journal of Machine Learning Research*, 24(89):4061–4157, 2023. arXiv:2108.08481.
- Z. Li, N. Kovachki, K. Azizzadenesheli, B. Liu, K. Bhattacharya, A. Stuart, and A. Anandkumar. Fourier Neural Operator for Parametric Partial Differential Equations. *ICLR*, 2021. arXiv:2010.08895.
- Z. Li, H. Zheng, N. Kovachki, D. Jin, H. Chen, B. Liu, K. Azizzadenesheli, and A. Anandkumar. Physics-Informed Neural Operator for Learning Partial Differential Equations. arXiv:2111.03794, 2023.
- L. Lu, P. Jin, and G. E. Karniadakis. Learning nonlinear operators via DeepONet based on the universal approximation theorem of operators. *Nature Machine Intelligence*, 3(3):218–229, 2021. arXiv:1910.03193.
- Y. Makara, Y. Tanaka, T. Matsubara, and T. Yaguchi. Symplectic Neural Operators for Learning Infinite Dimensional Hamiltonian Systems. arXiv:2605.15881, 2026.
- J. E. Marsden and T. S. Ratiu. *Introduction to Mechanics and Symmetry: A Basic Exposition of Classical Mechanical Systems*. Texts in Applied Mathematics, Vol. 17, 2nd edition, Springer, 1999.
- H. C. Öttinger. *Beyond Equilibrium Thermodynamics*. Wiley-Interscience, Hoboken, NJ, 2005.
- E. Perez, F. Strub, H. de Vries, V. Dumoulin, and A. Courville. FiLM: Visual Reasoning with a General Conditioning Layer. *AAAI*, 2018. arXiv:1709.07871.

- M. Raissi, P. Perdikaris, and G. E. Karniadakis. Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations. *Journal of Computational Physics*, 378:686–707, 2019.
- J. Sulskis and S. Ravi. GENERIC-FNO: Embedding Energy Conservation and Entropy Production into Fourier Neural Operators. arXiv:2606.08343, 2026.
- L. N. Trefethen. *Spectral Methods in MATLAB*. SIAM, Philadelphia, 2000.
- A. J. Varghese, S. Liu, P. Chen, Y. Zhu, J. Darbon, and G. E. Karniadakis. PI-SONet: A Physics-Informed Symplectic Operator Network for Real-Time Optimal Control of Multi-Agent Systems. arXiv:2605.14332, 2026.
- A. Weinstein. Symplectic structures on Banach manifolds. *Bulletin of the American Mathematical Society*, 75:1040–1041, 1969.
- G. Wen, Z. Li, K. Azizzadenesheli, A. Anandkumar, and S. M. Benson. U-FNO—An enhanced Fourier neural operator-based deep-learning model for multiphase flow. *Advances in Water Resources*, 163:104180, 2022. arXiv:2109.03697.
- Z. Zhang, Y. Shin, and G. E. Karniadakis. GFINNs: GENERIC formalism informed neural networks for deterministic and stochastic dynamical systems. *Philosophical Transactions of the Royal Society A*, 380(2229):20210207, 2022.
- J. Zhang, Y. Wang, and G. Lin. LegONet: Plug-and-Play Structure-Preserving Neural Operator Blocks for Compositional PDE Learning. arXiv:2603.07882, 2026.
- Y. D. Zhong, B. Dey, and A. Chakraborty. Symplectic ODE-Net: Learning Hamiltonian Dynamics with Control. *ICLR*, 2020. arXiv:1909.12077.

A Retraction of a previously asserted guarantee

The CKINO operator was originally documented with an exact symplecticity theorem. We retract it, and record the error because it is instructive.

The withdrawn argument ran: each Störmer–Verlet substep has a Jacobian that is triangular with unit diagonal blocks, hence determinant 1, hence preserves ω . The step from “determinant 1” to “preserves ω ” is invalid. By Proposition 2, $\det D\Phi = 1$ holds for every shear, while preservation of ω_W additionally requires $WA = A^\top W$. The operator’s kernel (5) uses independent φ_r, ψ_r and an unconstrained channel mix, so A is not self-adjoint; the measured defect is ≈ 1.4 in *both* candidate forms (Figure 5). A corollary energy-drift bound, whose proof invoked the modified-Hamiltonian machinery of backward error analysis (Hairer et al., 2006), is retracted with it: that machinery requires symplecticity, and volume preservation alone does not imply the existence of a modified Hamiltonian.

A second, subtler retraction concerns the comparison itself. An earlier draft reported that exact symplecticity was a measurable cost, on the basis of a defect table computed against a single inner product. Scoring both variants against both forms (Figure 5) shows the “non-symplectic” control was exactly symplectic in the uniform form all along, and reverses the conclusion.

B The weighted adjoint in the low-rank basis

With $\alpha_r = \langle \psi_r, v \rangle_W$ and $(Kv)(x) = \sum_r \varphi_r(x) M_r \alpha_r$,

$$\langle Kv, g \rangle_W = \sum_r (M_r \alpha_r)^\top \langle \varphi_r, g \rangle_W = \sum_r \alpha_r^\top (M_r^\top \langle \varphi_r, g \rangle_W) = \left\langle v, \sum_r \psi_r M_r^\top \langle \varphi_r, g \rangle_W \right\rangle_W,$$

which is exactly $\langle v, K^*g \rangle_W$. Hence the weighted adjoint is obtained by exchanging φ and ψ and transposing the channel mixing; no inverse of W is formed and the cost remains $O(RN^d)$. In d dimensions the weight is the outer product of the per-axis weights and remains diagonal, so both this identity and Theorem 5 carry over unchanged.